

Unit - 2

SQL, Intermediate SQL and Relational Database Design:

SQL: Overview of SQL, DDL, integrity Constraints, DML, Set Operations, null values, aggregate functions, Sub-queries

Intermediate SQL : Joins, Views, Indexes, Abstract Data type

* SQL - Structured Query Language

- Developed by IBM in the 1970's, Originally called Sequel
- The Standard relational-database language
- SQL is used to interact with a database to manage and retrieve data
- SQL is a non-procedural language

Data types

- Integer • Float • Char • Varchar2 • Number • Date

1) Integer : Used to store whole numbers (without decimal points)

Ex: 10, -50, 1000

2) Float : Used to store decimal numbers (Floating-point numbers)

Ex: 3.14, -0.75, 100.25

3) Char : • Stored fixed-length character strings

• If the defined length is 5, it always stores 5 characters
(even if fewer characters are used, spaces are padded)

Ex: CHAR(5) stores 'ABC--'

4) Varchar2 : Variable-length character data type

• Unlike CHAR, it does not pad spaces if the value is shorter than the defined length

Ex: VARCHAR2(10) can store 'Hello' (takes only 5 bytes, not 10).

5] NUMBER : Used in Oracle database to store both integers and floating point numbers

Ex: Number(5) → Can store up to 5-digit integers (99999)

Number(7,2) → Can store up to 7 digit with 2 decimal places (12345.67)

6] DATE : Stores date and time values.

Default Format in SQL : 'YYYY-MM-DD'

Ex: '2025-03-01' (March 1, 2025)

* DML is Data Manipulation Language Statements

Example: Select - Retrieves data from a table

Insert - Insert data into a table

Update - Modifies existing records

Delete - Removes records from table

* DDL is Data Definition Language

• Is used to define, modify and delete database structures

Create - Creates a new database object

Alter - Modifies an existing database object

Drop - Deletes an existing database object

Truncate - Removes all records from a table but keeps its structure

Rename - Renames a table or column

* DCL (Data Control Language)

• Is used to control access and permissions in a database

• It is mainly used for security purposes, allowing to grant or revoke privileges to users

Grant - Gives specific privileges to a user

Revoke - Removes specific privileges from a user

Commit - Save work done

Rollback - Restore database to original since the last commit

SQL - Create Table

Syntax: CREATE TABLE table-name {
 Column1 datatype Constraints,
 Column2 datatype Constraints,
 ...
};

Ex: Employee Table CREATE TABLE employees {
 id INT Primary key,
 name Nvarchar(20),
 Salary Decimal (10, 2),
};

ALTER TABLE

1. Add a New Column

Syntax: ALTER TABLE table-name
ADD Column-name datatype Constraints;

Ex: ALTER TABLE employees
ADD department NARCHAR(50);

2. Modify an Existing Column

Syntax: ALTER TABLE table-name {
 MODIFY Column-name newdatatype);

Ex: ALTER TABLE employees {
 MODIFY Salary DECIMAL (12, 2);

3. Rename a Column

Syntax: ALTER TABLE table-name
RENAME COLUMN Old-Column-name To new-Column-name;

Ex: ALTER TABLE employees
RENAME Column department To dept-name)

4. Drop (Delete) a column

Syntax: ALTER TABLE table-name
DROP COLUMN Column-name;

Ex: ALTER TABLE employees
DROP Column dept-name; } Note: All data in the column
will be lost permanently

* Difference Between TRUNCATE and DELETE

Both Truncate and Delete are used to remove records from a table, but they work differently

1. Delete Statement

- used to remove specific records (or all records) from a table.
- Rollback is possible (Delete is a DML Command)
- can include a WHERE Condition to delete selected rows

2. Truncate Statement

- Removes all rows from a table without a WHERE condition
- faster than DELETE (Truncate is a DDL Command)
- Cannot be rolled back (Permanent action)

* GROUP BY

- The Group By Clause is used in a SELECT Statement to collect data across multiple records and group the results by one or more columns
- Sometimes it is required to get information not about each row, but about each group.

* Integrity Constraints in SQL

- Integrity Constraints are rules enforced on table columns to ensure the accuracy and consistency of data in a relational database.

Types of Integrity Constraints

1. Primary Key Constraint

- Ensures that each row in a table is uniquely identified
- A table can have only one primary key
- The column(s) must be unique and NOT NULL

2. Foreign Key Constraints

- Ensures referential integrity between two tables.
- The foreign key column must reference a primary key from another table.

3. Unique Constraint

- Ensures that all values in a column are different
- Unlike Primary key, a table can have multiple unique columns

4. NOT NULL Constraint

- Prevents NULL values from being entered in a column.

5. Check Constraint

- Ensures that a column meets a specific condition
- Can be used to restrict values entered into a column.

6. Default Constraint

- Assigns a default value to a column if no value is provided

* Set Operations in SQL *

- Set operations in SQL are used to combine results from multiple SELECT queries.

1. UNION

2. UNION ALL

3. INTERSECT

1. UNION

- Combines results from two queries and remove duplicates
- Both queries must have the same number of columns and compatible data types.

Syntax: SELECT Col-1, Col-2 From table1

UNION

Select Col-1, Col-2 From table2

Example: Select name from Students

UNION

Select name from teachers;

Returns a list of unique names from both tables.

2. UNION ALL

- Similar to UNION, but keeps duplicate values
- Faster than UNION since it does not remove duplicates

Syntax: Select Column-1 From table1

UNION ALL

Select Column-1 From table 2;

Example: Select name from Students

UNION ALL

Select name from teachers;

Returns all names, including duplicates.

3. Intersect

- Returns only common records present in both queries
- The columns must have the same number and data type

Syntax: Select Col-1 From table1

INTERSECT

Select Col-1 From table2;

Example: Select name from Students

INTERSECT

Select name from teachers;

Returns names present in both tables

* NULL Values in SQL

- NULL represents a missing, unknown, or undefined value in a database table.
- It is not the same as 0, an empty string ("") or a space
- NULL means missing or unknown data, not zero or empty
- Use ISNULL or IS NOT NULL to filter NULL values
- Aggregate functions ignore NULL values

* Aggregate Functions

Aggregate functions perform calculations on multiple rows and return a single result. These are commonly used with the GROUP BY clause.

1. Count() - Counts Rows : Counts the number of rows in a column (excluding NULL values)

Syntax: Select Count (col-name) from table-name;

Example: 1] Select Count (Student-id) from Students;
→ Counts the total number of student

2] Select Count(*) from Students;
→ Counts all rows, including those with NULL values

2. Sum() - Calculates Total : Returns the sum of values in a numeric column

Syntax: Select SUM (col-name) From table-name;

Example: Select sum (Salary) From employees;

returns the total sum of salaries

3. AVG() - Calculates Average
• returns the average (mean) of numeric values

Syntax: Select AVG(col-name) From table-name;

Example: Select AVG(salary) From employees;
→ returns the average salary of employees

4. MAX() - Finds Maximum Value

• returns the highest value in a column

Syntax: Select MAX(col-name) From table-name;

Example: Select MAX(salary) From employees;
→ returns the highest salary

5. MIN() - Finds Minimum Value (Same as MAX())

6. Group By - Aggregating

The Group By clause is used to group data based on a column and apply aggregate functions to each group

Ex: Find total salary per department

```
Select department , SUM(Salary)  
From employees  
Group By department;
```

→ Groups employees by department and calculates the total salary in each department

7. HAVING - Filtering Aggregate Results

Unlike WHERE, which filters individual rows, HAVING filters grouped data

Ex: Find departments where total salary > 50,000

Select department, sum(salary) As total-salary
FROM employees
Group By department
Having sum(salary) > 50000;

→ Displays only department where the total salary exceeds 50,000

Sub-queries

- A Subquery (also called a nested query) is a query inside another SQL query.

* Types of Subqueries

1] Single-row Subqueries:

→ Returns one value (used with =, >, <, >=, <=, <>)

Ex: Find employees earning more than the average salary

Select name, salary
From employees

WHERE salary > (Select avg(salary) From employees);

→ The Subquery calculates the average salary, and the other query fetches employees earning above it.

2] Multiple-Row Subqueries

→ A Subquery that returns multiple values (IN, ANY, ALL)

Ex: Find employees working in the same department as 'Ritu'

Select name, department
From employees

Where department IN (Select department From employees
WHERE name = 'Ritu');

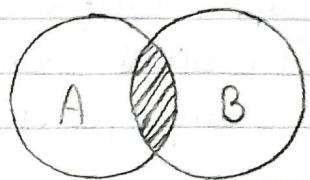
→ The subquery finds Ritu's department, and the outer query selects employees from the same department.

* SQL JOINS *

- SQL Joins are used to combine rows from two or more tables based on related column.
- Join operations take two tables and return another table as result.

INNER JOIN (Default Join)

→ Returns only the matching records from both tables



employees table

emp_id	name	dept_id
1	Ritu	101
2	Piyu	102
3	Aaru	103
4	Niku	104

departments table

dept_id	dept_name
101	HR
102	IT
103	Finance

Query: Get employee names with their department names

Select employees.name, departments.dept_name
FROM employees
INNER JOIN departments
ON employees.dept_id = departments.dept_id;

Output: name dept-name
 Ritu HR
 Piyu IT
 Aaru Finance

→ NIKU is not included because she has no matching department

2) LEFT JOIN



- returns all records from the left table and matching records from the right table.
- If no match is found, NULL is returned

Query: Get all employees and their departments (if any)

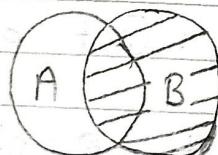
```
Select employees.name, departments.dept-name  
from employees  
LEFT JOIN departments  
ON employees.dept-id = department.dept-id;
```

Output:

Name	Dept-name
Ritu	HR
Piyu	IT
Aeru	Finance
Niku	NULL

→ Niku appears with NULL because he has no matching department

3) RIGHT JOIN



- returns all records from the right table and matching records from the left table
- If no match is found, NULL is returned

Query: Get all department and their employees

```
Select employees.name, department.dept-name  
from employees  
RIGHT JOIN departments  
ON employees.dept-id = department.dept-id;
```

Name	Dept-name
Ritu	HR
Piyu	IT
Aeru	Finance
NULL	Marketing

4] FULL JOIN



- Returns all records from both tables.
If there's no match, NULL is returned.

Query: Get all employees and all departments.

Select employees.name, departments.dept-name

From employees

FULL JOIN departments

ON employees.dept-id = departments.dept-id;

Output:

	name	dept-name
	Ritu	HR
	Piyu	IT
	Aaru	Finance
	Niku	NULL
	NULL	Marketing

→ Both unmatched employees and unmatched departments appear with NULL

5] CROSS JOIN

- Returns the Cartesian product of both tables
- every row from the first table pairs with every row from the second table

Query: Get all employee-department combinations

Select employees.name, departments.dept-name

From employees

CROSS JOIN departments;

Output:

	name	dept-name
	Ritu	HR
	Ritu	IT
	Ritu	Finance

Same for all other name

6) Self Join

- A table joins with itself, treating it as two separate tables
- Ex: find employees and their managers (assuming manager exists in employees)

Select e1.name As Employee, e2.name As Manager
From employees e1
Left Join employees e2
ON e1.manager-id = e2.emp-id;

C Output:

Employee

Ritu

Piyu

Aaru

Manager

Piyu

Aaru

NULL

→ Aaru has no manager (NULL).

→

③

q1

Se

Fs

RII

??